

Hyper-parameter optimization for convolutional neural networks using a Gaussian Process and simple partial randomization applied to the MNIST dataset [5]

Steven Smith

P13167208


Thursday 18th May – 15:40

What is to come?

- Problem definition
- Summary of existing literature
- Proposed solution
- Methodology
- Experimental design
- Data collection
- Experimental results
- Conclusions
- Future work
- Lessons learned
- References

Problem Definition

- Convolutional Neural Network
 - Convolutional Layers
 - Feature maps
 - Pooling Layers
 - Reduce features using maximum or average pooling
 - Fully Connected Layers
 - Traditional neural network layers




Problem Definition (continued)

- Hyper-parameters
 - Number of layers and filters
 - Size of filters and stride when applying filters
 - Regularization, e.g. L1, L2
 - Max pooling kernel size and strides
 - Number of neuronal units
 - Dropout percentage
 - Learning rate and decay
 - Activation functions, e.g. RELU, Linear, tanh, softmax, sigmoid
 - Optimization functions, e.g. RMSProp, Stochastic Gradient Descent & ADAM
 - and many many more!
- Grid Search
 - Large search space
 - The experiments uses 1,555,200 possible combinations
- Random Search
 - Not guided

Problem Definition (continued)

- MNIST [5]
- Digitised handwritten digits 0 to 9
- 70000 Digits in total
 - 60000 train/validate
 - 10000 test
- Best CNN error rate is 0.23% (Accuracy 99.77%)



Summary of existing literature

- Sequential Model-based Global Optimization
 - Iteratively maximize a ‘cheap’ surrogate function using an acquisition function such as Expected Improvement
 - At the maximum is where the parameters should now be tested against the real ‘expensive’ function
- Gaussian Process Regression [2]
 - Updating Multi-Variate Gaussian
 - Bayesian - Prior (existing knowledge) to Posterior (prediction)
 - Maximise Expected Improvement
 $EI(x) = \text{current best min} - (\mu + 1.96\sigma)$
 - “All hyper-parameters are not relevant for each point” [1, pp.3]
 - Multiple GPs for each layer
 - Issues modelling categorical data such as activation function
 - Matern 5/2 Kernel
 - As opposed to Squared Exponential Kernel

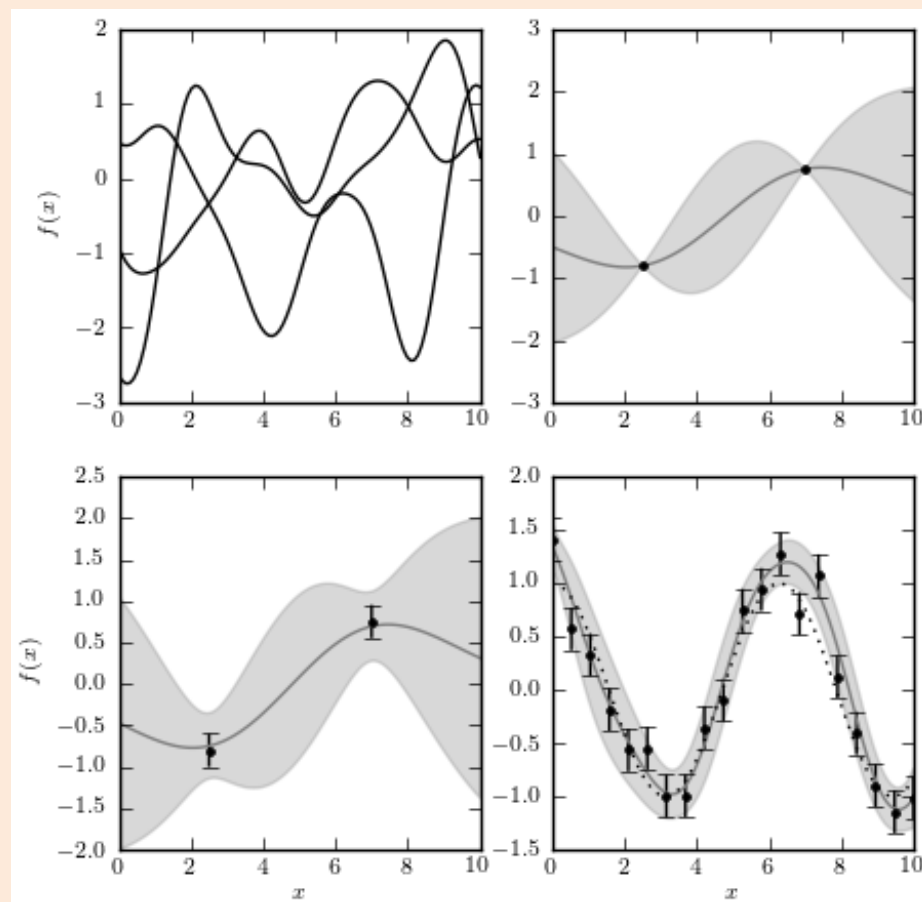


Image accessed on 09/05/2017 from http://www.astroml.org/book_figures/chapter8/fig_gp_example.html

Summary of existing literature (continued)

- Tree-Structured Parzen Estimator Approach [1]
 - Parameters modelled by separate distributions
 - Gather some random samples and split into best and worst observations for a parameter
 - Determine which group parameter sets are likely to belong to.
 - Expected Improvement $EI(x) = \frac{P(x \text{ is in best group})}{P(x \text{ is in worst group})}$
 - Parzen-window density estimator, non-parametric probability density function estimator
 - Tree based due to splitting on a hierarchy of parameters
- Transfer Learning
 - Use Gaussian Processes and learning built on multiple training datasets
 - One approach has access to the entire datasets [7]
 - The other uses statistics about the datasets already trained [8]
- Covariance Matrix Adaptation-Evolution Strategy
 - Preprint paper [10], uses an evolution based algorithm to determine the hyper-parameters
 - Found to have better results than both TPE and GP.
- Auto-Weka [11]
 - Implemented both GP and TPE based methods to determine which of 39 classifiers to use
 - Hyper-parameter optimization for those classifiers also.

Proposed solution

- In [1] authors suggest not all parameters need to be set, e.g. if layer 2 or 3 does not exist
 - The proposed approach uses the same hyper-parameters for different layers
- [1] also suggests that TPE outperformed GP due to “the exploration induced by the TPE’s lack of accuracy turned out to be a good heuristic for search” [1](pp.7)
 - The proposed approach adds partial randomization after GPR
- Gaussian Process Regression (GPR) optimizing using expected improvement of the log loss (or cross entropy) metric and random search
- Convolutional hyper-parameters vs fully connected hyper-parameters
- GPR to determine initial parameters then perturb some of the parameters for additional randomness.
 - This allows GPR to use all parameters to guide overall direction.
 - Only sets either convolutional or fully connected, the others are randomly generated.
 - Can this improve traversal of the search space?

Methodology

- Scope - 12 hyper-parameters converted to a (1 x 12) array for randomization or the Gaussian Regression Process
 - Number of conv. Layers, (1,3), $n = 1, 2, 3$
 - Number filters, (1,4), $2^{n+3} = 16, 32, 64, 128$
 - Filter size, (3,10), $n = 3, 4, 5, 6, 7, 8, 9, 10$
 - Stride, (1,5), $n = 1, 2, 3, 4, 5$
 - Regularizer, (1,2), $L_n = L1, L2$
 - Local Response Normalization, (0,1), 0 = off or 1 = on
 - Use Max Pooling, (0,1), 0 = off or 1 = on
 - Max Pool kernel size, (2,4), $n = 2, 3, 4$
 - Number of fully connected layers (1,3), $n = 1, 2, 3$
 - Number of neuronal units, (1,3), $2^{n+6} = 128, 256, 512$
 - Drop out keep probability (1,5), $0.1 \leq n \leq 0.4 = 0.5, 0.6, 0.7, 0.8, 0.9$
 - Learning rate (1,3), $10^{(-n)} = 0.1, 0.01, 0.001$
- For multiple layers, the same parameters are used
- Fixed parameters (taken from TFLearn defaults)
 - RELU for convolutional layer activation, Tanh for Fully connected layer activations
 - Weights are initialized using a truncated normal, $N(0, 0.02)$ and < 2 standard deviations from the mean
- Mini batches of size, 64
- All other network hyper-parameters are static between runs.
- Limit training to 20 epochs (50 epochs for the final-full tests)

Methodology (continued)

Four separate methods tested:

1. Random search
2. GPR without further randomization
3. GPR with 8 Convolutional parameters randomized
[1,2,3,4,5,6,7,8,1,2,3,1] by GPR
[5,6,7,8,1,2,3,4,1,2,3,1] further randomized
4. GPR with 3 Fully Connected Parameters randomized
[1,2,3,4,5,6,7,8,1,2,3,1] by GPR
[1,2,3,4,5,6,7,8,9,8,7,1] further randomized

Experimental design

- The four experiments in each batch
- Each batch has
 - Number of iterations – 20 or 50
 - Number of epochs – 20
 - Validation set train/validate ratio – fixed at 70/30
 - Groups of MNIST [5] digits to train on
 - 4,7
 - 4,7,9
 - 0, 4, 7 ,9
 - 0, 2, 4, 7 ,9
 - 0, 1, 2, 3, 4, 5, 6, 7, 8, 9
- Ten combinations in all (Training 1400 models, 350 per algorithm, for 20 epochs each)

Data collection


- Using Python, TFLearn, Tensorflow, NumPy, Sci-kit Learn, Matplotlib to train the models and generate data.
- Generate plots of accuracy and log loss to compare visually.
- Accuracy, Log Loss, Confusion Matrix for every trained model recorded in a log of each experiment.
- Processed by jupyter notebook to extract data
- Excel to create tables and graphs

Experimental results


	Maximum Accuracy			
	Random	Gaussian_EI	Gaussian_EI_RC	Gaussian_EI_RF
2n_20	0.995	0.996	0.979	0.954
2n_50	0.998	0.981	0.983	0.988
3n_20	0.984	0.986	0.985	0.341
3n_50	0.974	0.958	0.907	0.861
4n_20	0.953	0.978	0.988	0.469
4n_50	0.982	0.995	0.984	0.857
5n_20	0.965	0.386	0.204	0.593
5n_50	0.949	0.988	0.854	0.980
10n_20	0.952	0.963	0.982	0.880
10n_50	0.926	0.967	0.988	0.918
Max	0.998	0.996	0.988	0.988
Count	3	4	3	0

	Count of Accuracy > 0.6			
	Random	Gaussian_EI	Gaussian_EI_RC	Gaussian_EI_RF
2n_20	4	1	2	1
2n_50	11	4	4	9
3n_20	5	2	2	0
3n_50	4	6	1	3
4n_20	2	1	2	0
4n_50	6	16	5	4
5n_20	2	0	0	0
5n_50	3	12	3	3
10n_20	1	11	8	2
10n_50	4	35	26	9
Total	42	88	53	31
Count	5	5	1	0

Experimental Results (continued)



Experimental Results (continued)



Experimental Results (continued)

- Best parameters for [0,1,2,3,4,5,6,7,8,9] and 50 iterations
 - Accuracy 98.79% (State of the art CNN Accuracy 99.77%)
 - [3, 1, 10, 1, 2, 1, 1, 3, 1, 3, 5, 3]
 - 3 Convolutional layers
 - filters=16
 - size=10
 - stride=1
 - L2 regularizer
 - local response normalization layer
 - max pooling
 - kernel size=3
 - 1 Fully Connected Layer
 - 512 units
 - Drop out keep probability = 0.9
 - Learning rate = 0.001

Conclusions

- Random, Gaussian_EI and Gaussian_EI_RC all had best performance in 3, 4 and 3 tests respectively.
- Therefore Gaussian Process was better than Random in 7 out of 10 tests.
 - However 8 of the 12 parameters are randomized so is that more a random process?
- Gaussian RC performed best on the Full Dataset on both 20 and 50 iterations
- When the parameters are far from optimum the time to complete epochs can be 4/5 times longer than other iterations
- Would a better approach have been to perturb the parameters to a lesser degree?
 - Only perturbing some of the parameters
 - Not perturbing over the whole range but close to the Gaussian process generated hyper-parameters.

Future work

- Incremental initialisation
 - Start with two digit classes, use GPR to find good models
 - Use these to initialise GPR against three digits, then four digits ...
- Time limited test runs as opposed to epochs or number of mini batches
- Can the optimizing of the Convolutional Layers be treated independently of the Fully Connected layers and vice versa?
- Random weights for feature learning [6]
- Create repository of good configurations from other papers
- Test against CIFAR-10 and CIFAR-100 [3,4]
- Reinforced Learning optimizing hyper-parameters on 2 classes then 3 classes then 4 ...

Lessons learned

- Start small, start simple
 - Difficult to make a complex solution simpler
- Use relational database to store experiment logs and results
- Split the learning across all available GPUs/CPU's
- Don't close the program when you mean to minimise it! (Twice!)

References

- [1] Bergstra, J.S., Bardenet, R., Bengio, Y. and Kégl, B., 2011. Algorithms for hyper-parameter optimization. In *Advances in Neural Information Processing Systems* (pp. 2546-2554).
- [2] Snoek, J., Larochelle, H. and Adams, R.P., 2012. Practical bayesian optimization of machine learning algorithms. In *Advances in neural information processing systems* (pp. 2951-2959).
- [3] Thornton, C., Hutter, F., Hoos, H.H. and Leyton-Brown, K., 2013, August. Auto-WEKA: Combined selection and hyperparameter optimization of classification algorithms. In Proceedings of the 19th ACM SIGKDD international conference on Knowledge discovery and data mining (pp. 847-855). ACM.
- [4] CIFAR-10 and CIFAR-100 datasets, <https://www.cs.toronto.edu/~kriz/cifar.html> last accessed 09/05/2017
- [5] The MNIST DATABASE of handwritten digits, <http://yann.lecun.com/exdb/mnist/> last accessed 09/05/2017
- [6] Saxe, A., Koh, P.W., Chen, Z., Bhand, M., Suresh, B. and Ng, A.Y., 2011. On random weights and unsupervised feature learning. In *Proceedings of the 28th international conference on machine learning (ICML-11)* (pp. 1089-1096).
- [7] Bardenet, R., Brendel, M., Kégl, B. and Sebag, M., 2013, June. Collaborative hyperparameter tuning. In ICML (2) (pp. 199-207).
- [8] Yogatama, D. and Mann, G., 2014. Efficient transfer learning method for automatic hyperparameter tuning. *Transfer*, 1, p.1
- [9] Rasmussen, C.E. and Williams, C.K., 2006. Gaussian processes for machine learning. 2006. The MIT Press, Cambridge, MA, USA, 38.
- [10] Loshchilov, I. and Hutter, F., 2016. CMA-ES for Hyperparameter Optimization of Deep Neural Networks. arXiv preprint arXiv:1604.07269.
- [11] Thornton, C., Hutter, F., Hoos, H.H. and Leyton-Brown, K., 2013, August. Auto-WEKA: Combined selection and hyperparameter optimization of classification algorithms. In Proceedings of the 19th ACM SIGKDD international conference on Knowledge discovery and data mining (pp. 847-855). ACM.